



# VISUM2020 PROJECT

# Fish Detection in Underwater Images

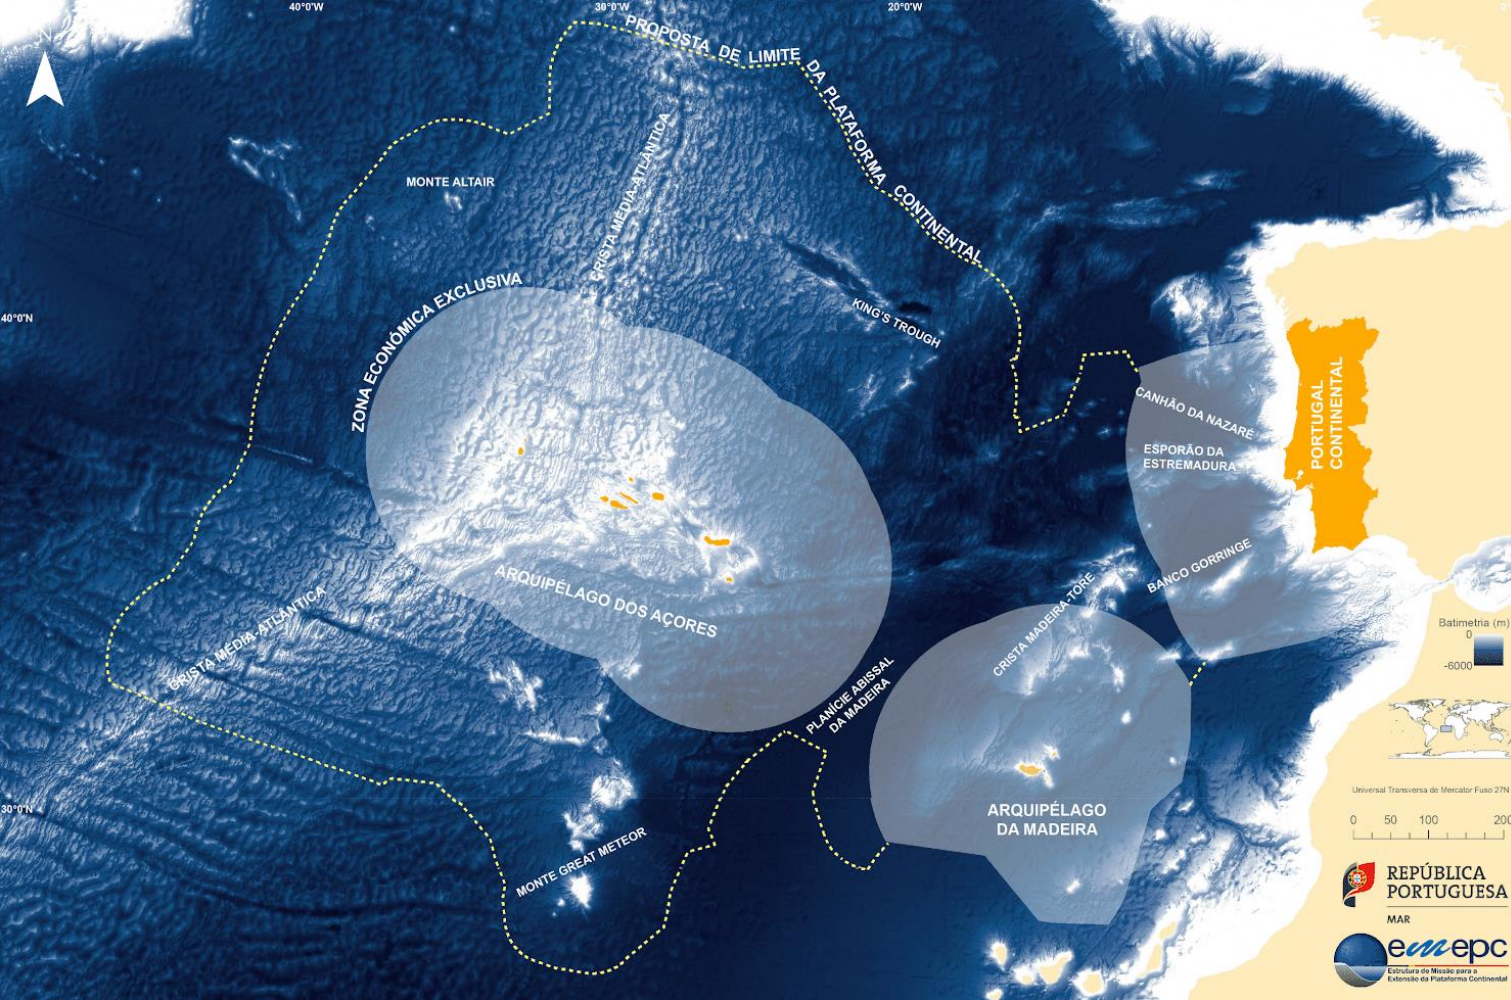


In collaboration with:



**abyssal**





# EMEPC

- Dataset was obtained by EMEPC
- They explore the Portuguese continental shelf
- Scientific work: geology and biology
- Close to 1000 hours of operation



# Why is it important to EMEPC?

Marine Biologists need to see the video to:

- Count fish
- Classify species

**Less than 1% of the video frames shows fish!**



# Why is it important to Abyssal?

- Offshore aquaculture
- Environmental studies
  - Before building a new offshore field (*i.e.* offshore wind farm)
  - Before decommissioning an old field
- Keep collaboration with EMEPC

# Resources

Each team will have access to a Google Cloud machine:

- NVIDIA Tesla K80
- 8 vCPU
- 30 GB RAM
- 60 GB Hard Disk



Google Cloud



# Using your machine

## Connecting

```
visum@instance-1: ~  
C:\Users\edux3\.ssh>ssh visum@35.243.226.67  
visum@35.243.226.67's password:  
Welcome to Ubuntu 18.04.2 LTS (GNU/Linux 4.15.0-1036-gcp x86_64)  
  
* Documentation:  https://help.ubuntu.com  
* Management:    https://landscape.canonical.com  
* Support:       https://ubuntu.com/advantage  
  
System information as of Fri Jul  5 16:24:44 UTC 2019  
  
System load:  0.0           Processes:            149  
Usage of /:   18.8% of 57.98GB Users logged in:       0  
Memory usage: 1%           IP address for ens5: 10.142.0.20  
Swap usage:   0%  
  
* MicroK8s 1.15 is out! It has already been installed on more  
  than 14 different distros. Guess which ones?  
  
  https://snapcraft.io/microk8s  
  
36 packages can be updated.  
0 updates are security updates.  
  
Last login: Fri Jul  5 13:32:24 2019 from 192.136.49.35  
visum@instance-1:~$ ls  
other_stuff  stuff  test.py  train.py  
visum@instance-1:~$
```

- `ssh visum@<ip_address>`
  - User: visum
  - Password: \*\*\*\*\*
- **All the users share the same account;**
- **There is no GUI.**

# Using your machine

## *Editing files*

```
visum@instance-1: ~  
C:\Users\edux3\.ssh>ssh visum@35.243.226.67  
visum@35.243.226.67's password:  
Welcome to Ubuntu 18.04.2 LTS (GNU/Linux 4.15.0-1036-gcp x86_64)  
  
* Documentation:  https://help.ubuntu.com  
* Management:    https://landscape.canonical.com  
* Support:       https://ubuntu.com/advantage  
  
System information as of Fri Jul  5 16:24:44 UTC 2019  
  
System load:  0.0           Processes:            149  
Usage of /:   18.8% of 57.98GB Users logged in:       0  
Memory usage: 1%           IP address for ens5: 10.142.0.20  
Swap usage:   0%  
  
* MicroK8s 1.15 is out! It has already been installed on more  
  than 14 different distros. Guess which ones?  
  
  https://snapcraft.io/microk8s  
  
36 packages can be updated.  
0 updates are security updates.  
  
Last login: Fri Jul  5 13:32:24 2019 from 192.136.49.35  
visum@instance-1:~$ ls  
other_stuff  stuff  test.py  train.py  
visum@instance-1:~$
```

- You may map your machine's /home to a folder in your computer.
  - This will allow you to use fancy text editors, view images and plot results

```
sshfs visum@:/home/visum ~Desktop/visum
```



# The Dataset

## TRAIN

- 84 sequences of 60 frames each;
- Some images contain one or more fish;
- Some images do not contain any fish.

`/home/master/dataset/train/seq000/img0.jpg`

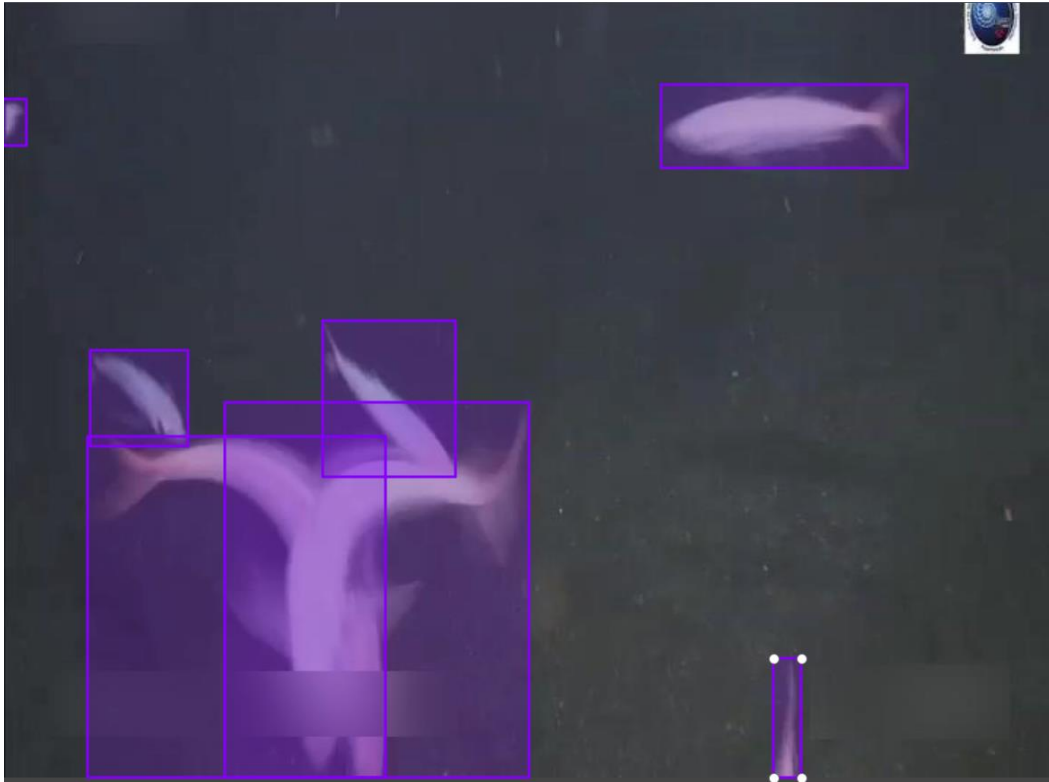
## TEST (NO ACCESS)

- 30 sequences of 60 frames each;
- Some images contain one or more fish;
- Some images do not contain any fish;
- Daily test: 10 sequences of 60 frames each.

# Annotations

## *CSV File*

1. sequence; frame; [(x\_min, y\_min, x\_max, y\_max), (x\_min, y\_min, x\_max, y\_max),...]
2. sequence; frame;



Can be found in:

`/home/master/dataset/train/labels.csv`

# Inference

- Test file path: /home/visum/test.py
- Should infer: /home/master/dataset/test
- Should generate a **valid** predictions.csv at /home/visum/predictions.csv

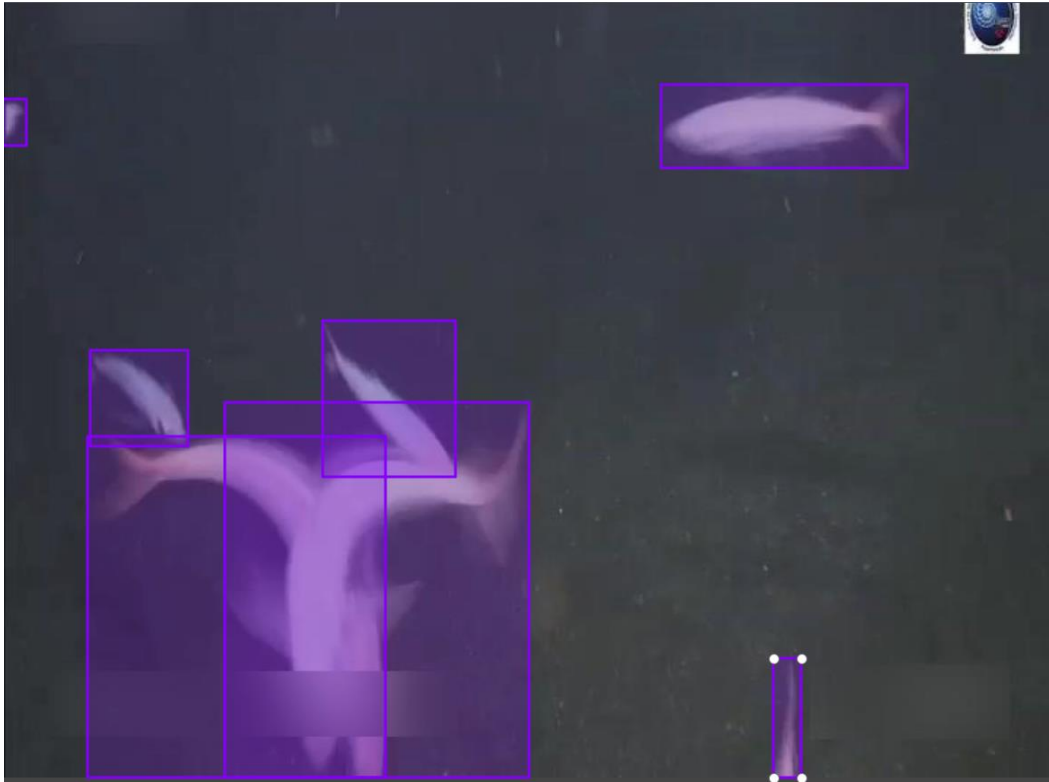
Your home/visum/ will be copied to our **server** and we will run inference and scoring **there**

```
visum@instance-1: ~  
visum@instance-1:~$ ls  
other_stuff  stuff  test.py  train.py  
visum@instance-1:~$ _
```

# Predictions

## *CSV File*

1. `sequence; frame; [(x_min, y_min, x_max, y_max)]; confidence`



**Must be saved in:**

`/home/visum/predictions.csv`



*“How do I test if my submission is valid?”*

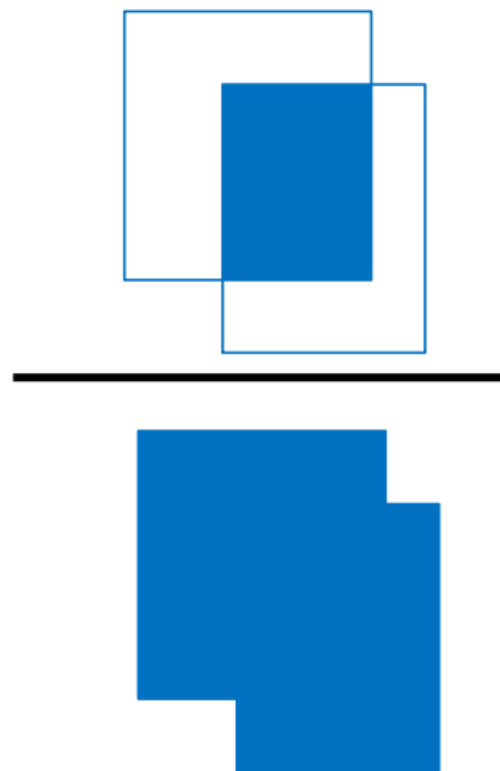
- Run:
  - `cd /home/visum/`
  - `rm predictions.csv`
  - `python test.py`
  - `python evaluate.py`

# Validating Submissions

# Evaluation

*IOU*

$$IOU = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

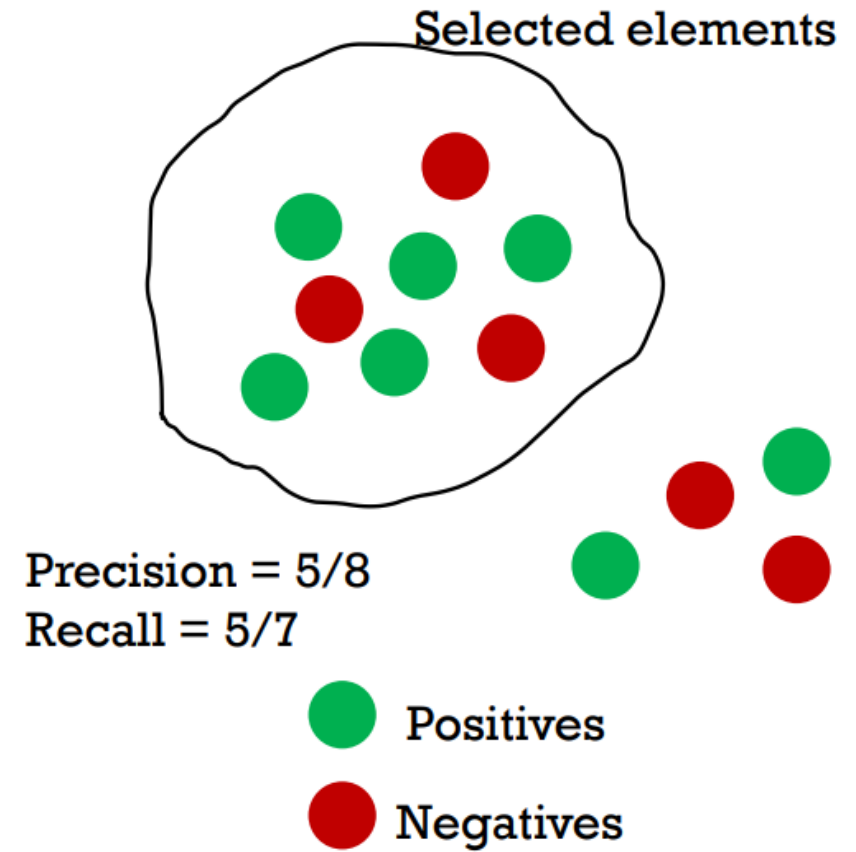


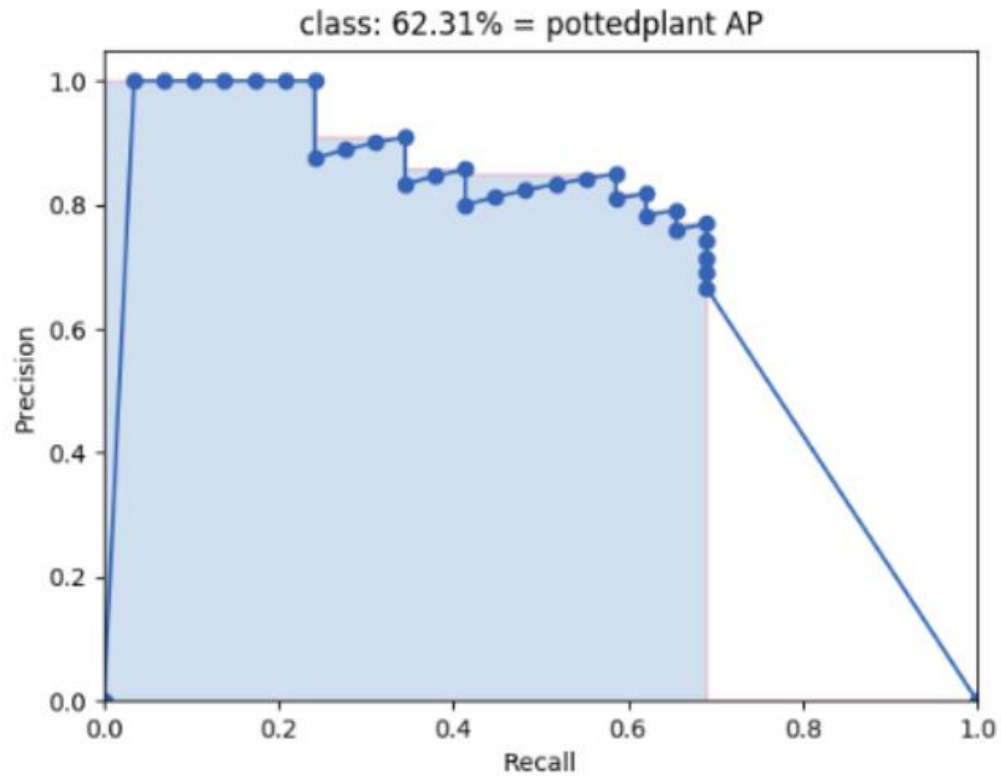
# Evaluation

## *Precision and Recall*

$$\textit{Precision} = \frac{TPs}{\textit{Total detections}}$$

$$\textit{Recall} = \frac{TPs}{\textit{Total objects}}$$





<https://github.com/cartucho/map>

# Evaluation

## *Average Precision*

**AP@[0.5:0.95]**

- Averaged over different IoU thresholds



# Daily, Final and Current Leaderboard

- Should run every day at around 12am
- Additional run whenever the team wants!
- Runs on a subset of the test set

## Daily Leaderboard (July 3)

	Team	AP@[0.5:0.95]
1	test_team_ws	0.1238
2	test_team_mf	0.1160
3	<i>BASELINE</i>	0.0707

# Competition Duration (*Purple blocks*)

\* All times are UTC/GMT+1 (Western European Summer Time)

	July, 03 (Fri)	July, 06 (Mon)	July, 07 (Tue)	July, 08 (Wed)	July, 09 (Thu)	July, 10 (Fri)
09:00		ML & CV BASICS (Jaime Cardoso) 08:30* - 10:00*				
10:00		Coffee Break	EXPLAINABLE AI (Mauricio Reyes) 09:00* - 11:30*	OPTIMAL TRANSPORT IN COMPUTER VISION (Nicolas Courty) 09:00* - 11:30*	INFORMATION SECURITY (Marta Gomez-Barrero) 09:00* - 11:30*	
11:00		MACHINE LEARNING ON KUBERNETES (Markus Bauer) 10:30* - 12:30*				Industry Session 11:00* - 12:30*
12:00			Lunch 11:30* - 13:30*	Coffee Break + AI shot	Coffee Break + AI shot	
13:00		Lunch 12:30* - 14:00*		Lunch 12:00* - 14:00*	Lunch 12:00* - 14:00*	Lunch 12:30* - 14:00*
14:00			DEEP LEARNING & KERNEL MACHINES (Johan Suykens) 13:30* - 16:00*	INESC Presentation		"AI: More than Human" Debate Session 14:00* - 15:30*
15:00	Opening Session 15:00* - 16:00*	ACTION RECOGNITION IN VIDEO (Pascal Mettes) 14:00* - 16:30*				
16:00	Project Presentation 16:00* - 17:00*		Coffee Break + AI shot			Project Presentations 15:30* - 17:00*
17:00	Welcome Session 17:00* - 18:00*	Coffee Break + AI shot		Project 14:30* - 20:00*	Project 14:00* - 20:00*	Closing Session 17:00* - 18:00*
18:00		Project 17:00* - 20:00*	Project 16:30* - 20:00*			
19:00						

You must use python3 as a programming language;

Your machine has some software installed: CUDA, cuDNN, Numpy, Open-CV, TensorFlow, PyTorch, Scikit-Learn;

You may not install software packages other than python packages;

Please, do not use virtual environments (e.g. conda);

We recommend you to not use Jupyter notebooks;

Your script (test.py) should run in acceptable time (less than one hour).

# Development

## Teams are eligible to win the project competition if:

1. At the end of the competition they have a valid submission;
2. They have been chosen for an oral presentation at the last the day of the summer school;

The winner will be determined by a jury composed of both VISUM and Abyssal members. The main criterion will be the highest average precision. However, creativity, novelty, and the ability to communicate ideas will also be considered in the final decision.

# Winning the project competition



# The Prize

By participating in our competition, you can win:

- A 100€ Amazon Voucher
- Free registration to VISUM 2021



# Finding help

If you have any **questions** regarding the project, relating to the rules, difficulties in understanding code or in setting up the google cloud machines, please find one member of the project staff:

- Eduardo - [eduardo.m.castro@inesctec.pt](mailto:eduardo.m.castro@inesctec.pt)
- Isabel - [icrto@fe.up.pt](mailto:icrto@fe.up.pt)
- João - [joao.t.pinto@inesctec.pt](mailto:joao.t.pinto@inesctec.pt)
- Mafalda - [mafalda.falcao@fe.up.pt](mailto:mafalda.falcao@fe.up.pt)
- Ricardo - [ricardo.j.araujo@inesctec.pt](mailto:ricardo.j.araujo@inesctec.pt)
- Wilson - [wilson.j.silva@inesctec.pt](mailto:wilson.j.silva@inesctec.pt)

## Network Goal: Object Detection

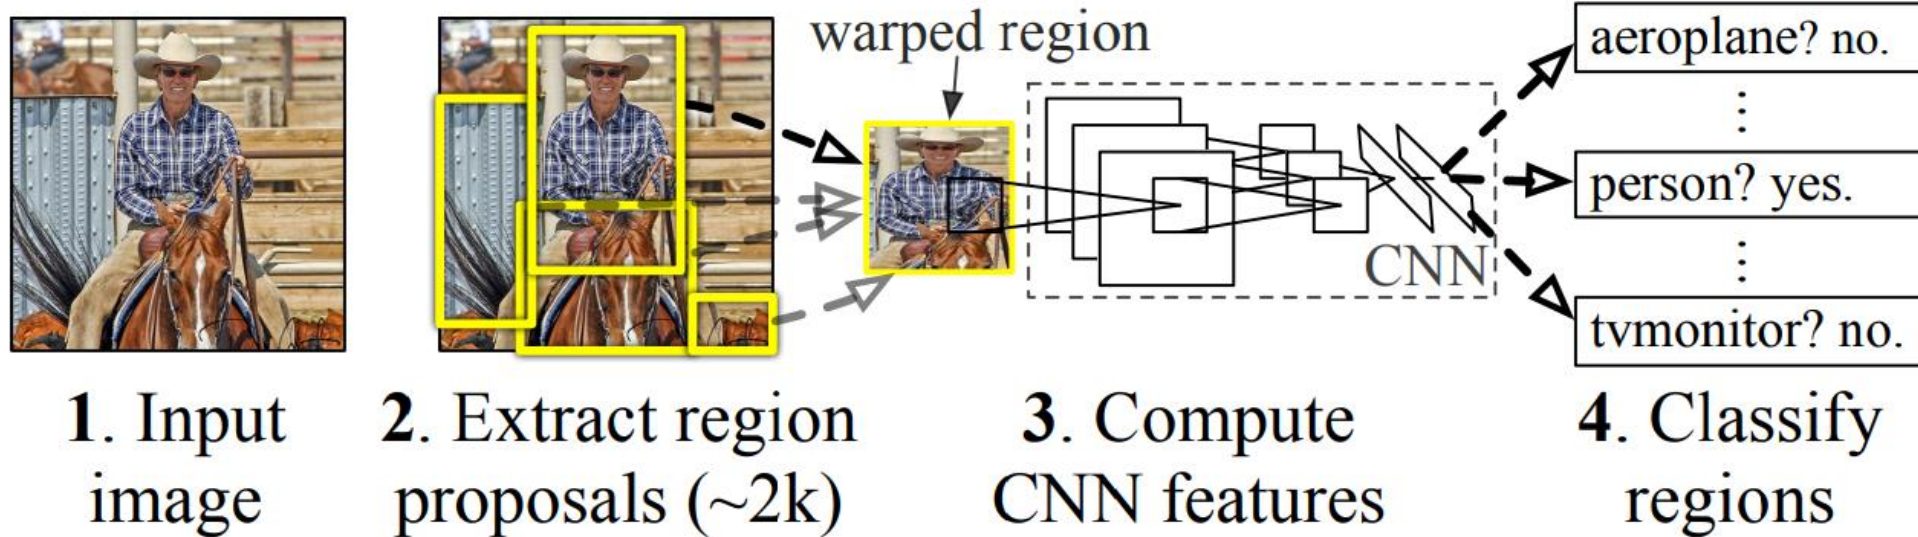
### State-of-the-art Architectures:

- YOLO (<https://arxiv.org/pdf/1506.02640.pdf>);
- RetinaNet (<https://arxiv.org/abs/1708.02002.pdf>);
- **Faster R-CNN** (<https://arxiv.org/pdf/1506.01497.pdf>).

# Developing a Baseline

# The R-CNN Architecture

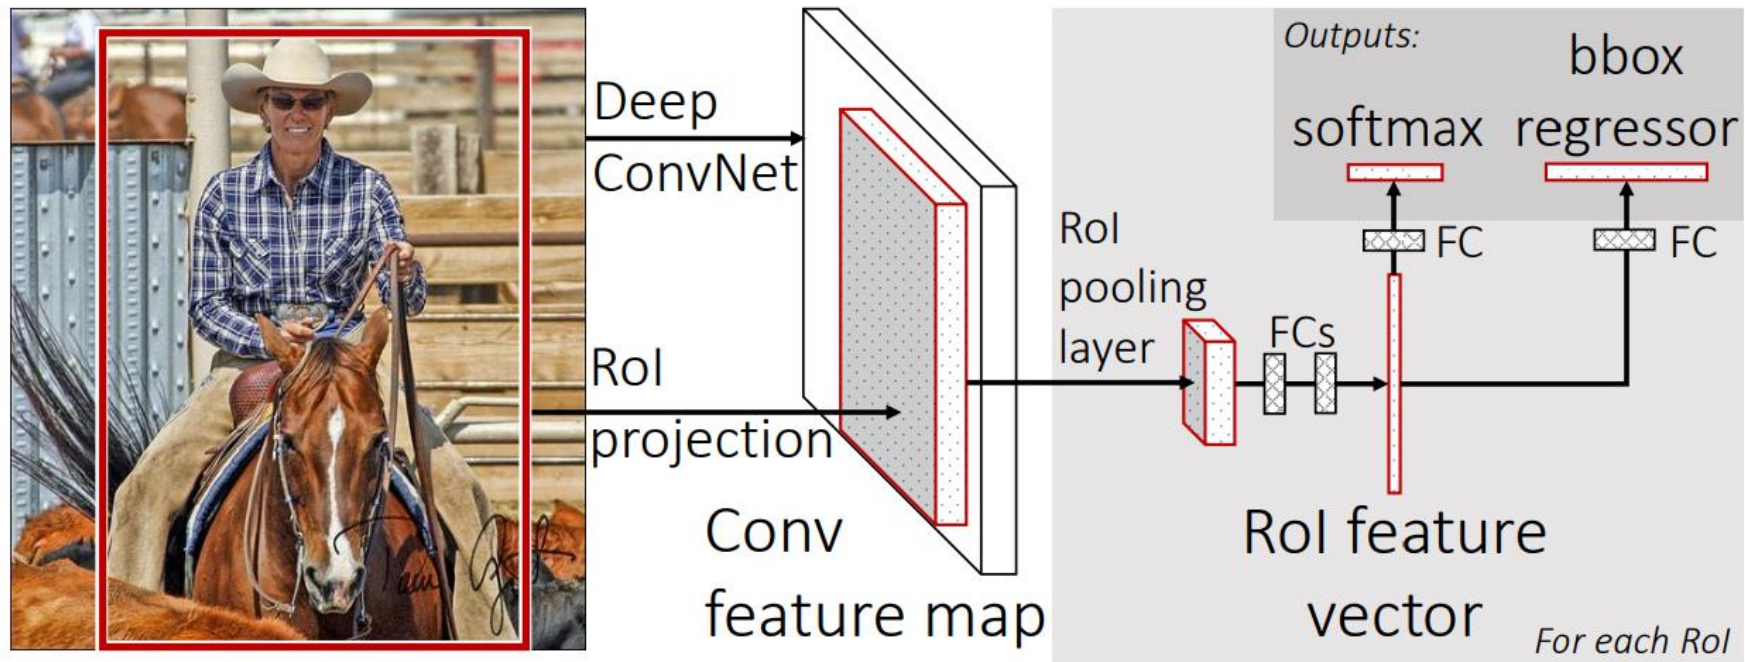
## R-CNN: *Regions with CNN features*



<https://arxiv.org/pdf/1311.2524.pdf>



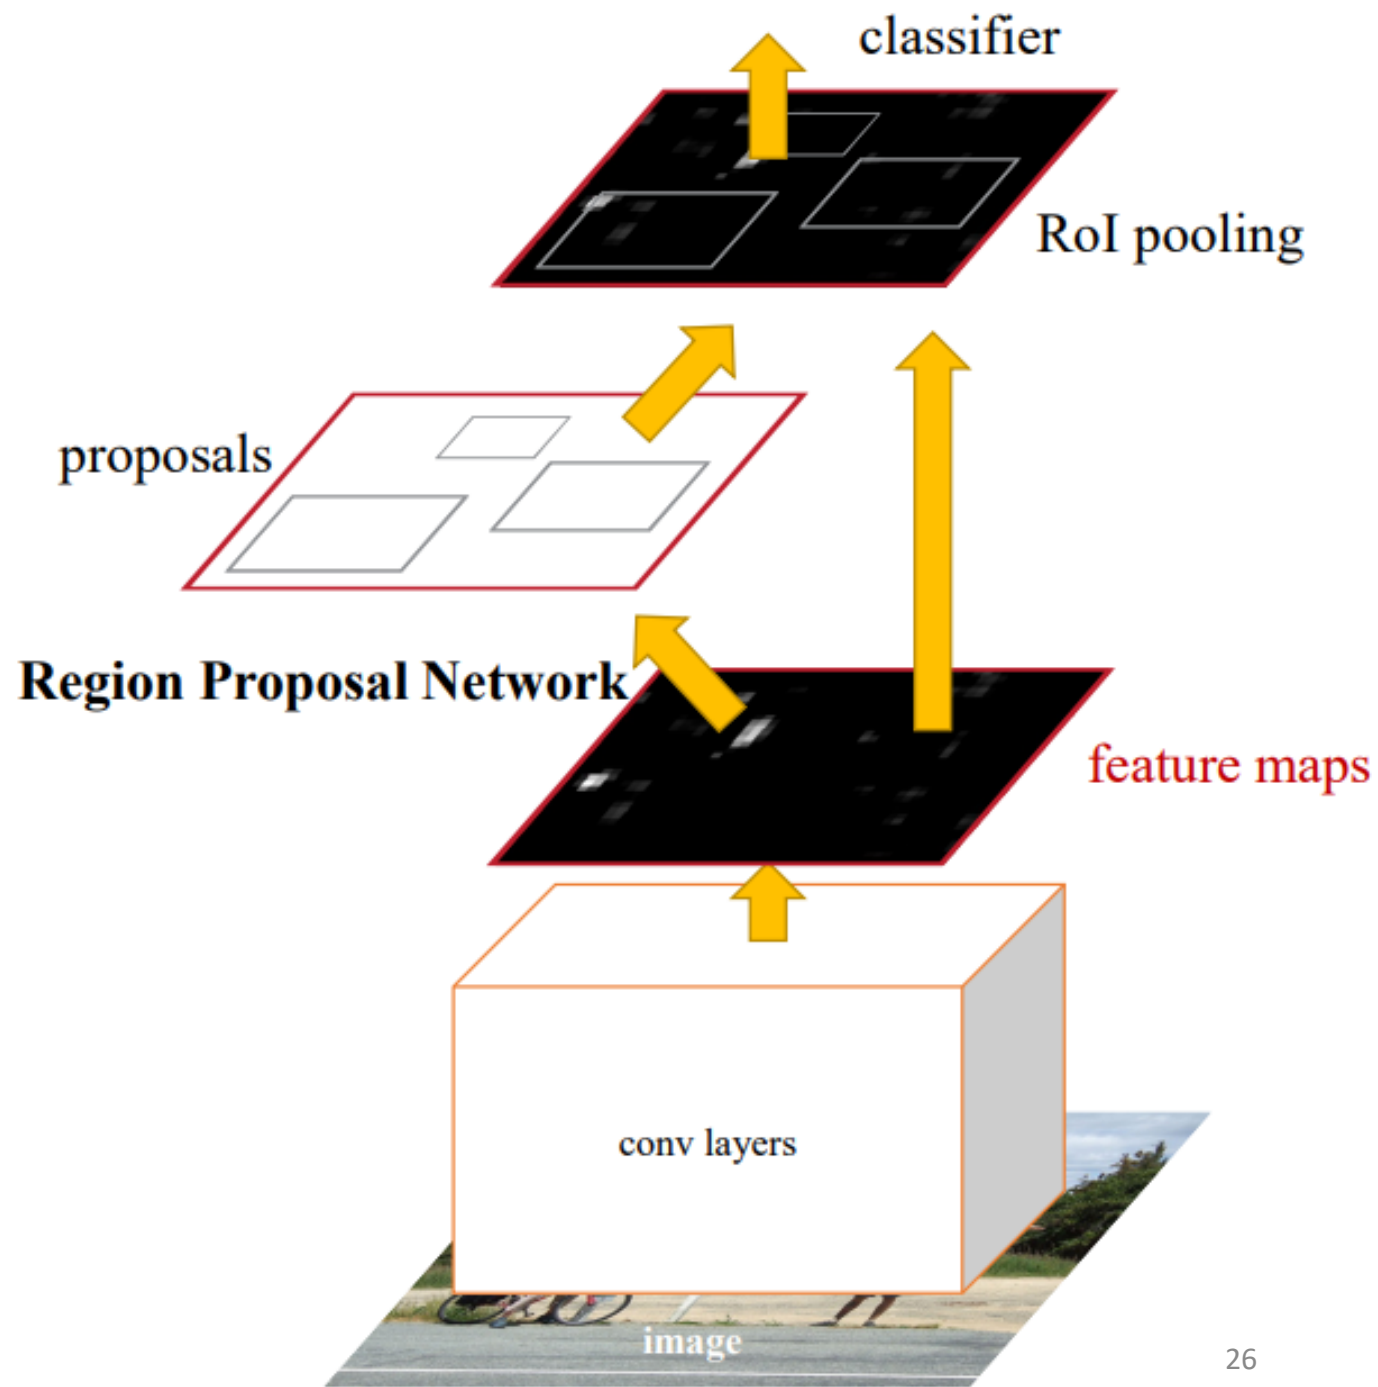
# The Fast R-CNN Architecture



<https://arxiv.org/pdf/1504.08083.pdf>

# The Faster R-CNN

*The Baseline*



## Implementation:

- PyTorch (*torchvision*)

<https://pytorch.org/docs/stable/torchvision/models.html#object-detection-instance-segmentation-and-person-keypoint-detection>

- Adapted from Mask R-CNN tutorial

[https://colab.research.google.com/github/pytorch/vision/blob/master/tutorial/tutorials/torchvision\\_finetuning\\_instance\\_segmentation.ipynb](https://colab.research.google.com/github/pytorch/vision/blob/master/tutorial/tutorials/torchvision_finetuning_instance_segmentation.ipynb)

# The Baseline Solution

# The Baseline Solution

## Files:

- `train.py` - Train and save your model;
- `test.py` - Load your trained model, post-process it and make predictions;
- `plot_results.py` - Visualize your predictions;
- `evaluate.py` - Check your scores!

# The Baseline Solution

## train.py

```
# load a pre-trained model for classification and return
# only the features
backbone = torchvision.models.mobilenet_v2(pretrained=True).features
# FasterRCNN needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each featuremap
# could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(
    sizes=((32, 64, 128, 256)), aspect_ratios=((0.5, 1.0, 2.0),)
)

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# OrderedDict[Tensor], and in featmap_names you can choose which
# feature maps to use
roi_pooler = torchvision.ops.MultiScaleRoIAlign(
    featmap_names=["0"], output_size=7, sampling_ratio=2
)

# put the pieces together inside a FasterRCNN model
# one class for fish, other for the background
model = FasterRCNN(
    backbone,
    num_classes=2,
    rpn_anchor_generator=anchor_generator,
    box_roi_pool=roi_pooler,
    min_size=300, max_size=300
)
```

## Model Definition



# The Baseline Solution

## train.py

```
# use our dataset and defined transformations
dataset = Dataset(DATA_DIR, transforms=get_transform(train=True))
dataset_val = Dataset(DATA_DIR, transforms=get_transform(train=False))

# split the dataset into train and validation sets
torch.manual_seed(1)
indices = torch.randperm(len(dataset)).tolist()

dataset_sub = torch.utils.data.Subset(dataset, indices[:-500])
dataset_val_sub = torch.utils.data.Subset(dataset_val, indices[-500:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset_sub, batch_size=6, shuffle=True, num_workers=4, collate_fn=utils.collate_fn
)

data_loader_val = torch.utils.data.DataLoader(
    dataset_val_sub, batch_size=6, shuffle=False, num_workers=4, collate_fn=utils.collate_fn
)
```

Split data into training  
and validation.

# The Baseline Solution

## train.py

```
device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
print(device)

model.to(device)

# define an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005, momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=15, gamma=0.1)

num_epochs = 20

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    epoch_loss = train_one_epoch(model, optimizer, data_loader,
                                  device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the validation dataset
    evaluator = evaluate(model, data_loader_val, dataset_val, device)

    torch.save(model, SAVE_MODEL)
```

Define optimizer and  
train your model.  
Finally, save it!

# The Baseline Solution

test.py

```
# Load dataset
dataset_test = Test_Dataset(DATA_DIR, transforms=get_test_transform())

device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# Load model
model = torch.load(SAVED_MODEL)
model.to(device)
```

Load your data and your model.

# The Baseline Solution

## test.py

```
predictions = list()

for ii, (img, seq, frame) in enumerate(dataset_test):
    if ii%50 == 0:
        print("Processed %d / %d images" % (ii, len(dataset_test)))

    # put the model in evaluation mode
    model.eval()
    with torch.no_grad():
        prediction = model([img.to(device)])

    boxes = prediction[0]['boxes'].cpu()
    scores = prediction[0]['scores'].cpu()

    nms_indices = nms(boxes, scores, NMS_THRESHOLD)

    nms_boxes = boxes[nms_indices].tolist()
    nms_scores = scores[nms_indices].tolist()

    # if there are no detections there is no need to include that entry in the predictions
    if len(nms_boxes) > 0:
        for bb, score in zip(nms_boxes, nms_scores):
            predictions.append([seq, frame, list(bb), score])

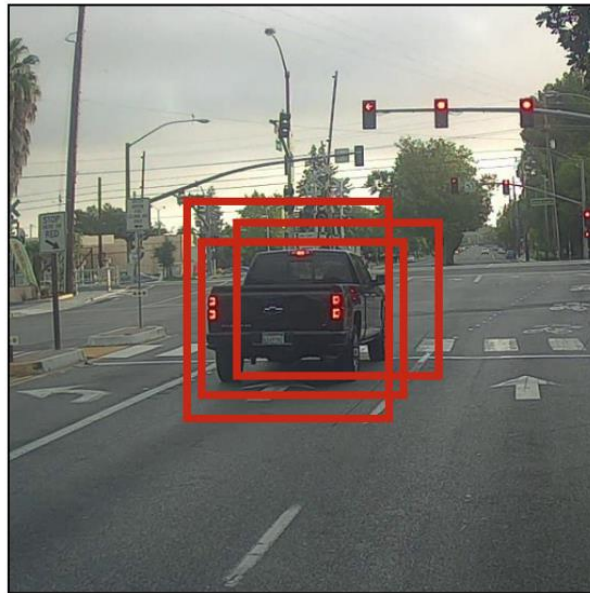
with open('predictions.csv', 'w', newline='') as file:
    writer = csv.writer(file, delimiter=";")
    writer.writerow(['seq', 'frame', 'label', 'score'])
    writer.writerows(predictions)
```

Make predictions, post-process them,  
and finally, save them in a csv file.

# The Baseline Solution

test.py

Before non-max suppression



Non-Max  
Suppression



After non-max suppression



Source: <https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

# The Baseline Solution

## plot\_results.py

```
def plot_frame(pred, gt, seq, frame, path):

    img = load_img(seq, frame)
    if(img is None):
        return

    draw = ImageDraw.Draw(img)

    pred_bb = load_bb(pred, seq, frame)
    gt_bb = load_bb(gt, seq, frame).iloc[0]
    scores = get_score(pred, seq, frame)

    # plot predicted bounding boxes (in blue)
    for bb, score in zip(pred_bb, scores):
        bb = ast.literal_eval(bb)
        xmin, ymin, xmax, ymax = bb[0], bb[1], bb[2], bb[3]
        draw.rectangle([(xmin, ymin), (xmax, ymax)], outline=(0, 0, 255), width=3)
        draw.text((xmin, ymin - 20), str(np.round(score, 3)), fill=(0, 0, 255), font=FONT)

    # plot ground truth bounding boxes (in green)
    if(gt_bb):
        for xmin, ymin, xmax, ymax in gt_bb:
            draw.rectangle([(xmin, ymin), (xmax, ymax)], outline=(0, 255, 0), width=3)

    # save image and create necessary paths
    img_path = os.path.join(path, 'seq' + seq)
    if(not os.path.exists(img_path)):
        os.makedirs(img_path)

    img_name = os.path.join(img_path, 'img' + frame + '.jpg')
    img.save(img_name)

def plot_sequence(pred, gt, seq, path):

    frames = gt[gt['seq'] == seq]['frame']
    for frame in frames:
        plot_frame(pred, gt, seq, frame, path)
```

Save a frame/sequence with its ground truth and predicted bounding boxes.

# The Baseline Solution

## plot\_results.py

```
if __name__ == '__main__':  
  
    # Load predictions  
    pred = pd.read_csv(PRED_PATH, delimiter=';', dtype={'seq': str, 'frame': str})  
  
    # Load ground truth  
    gt = pd.read_csv(GT_PATH, delimiter=';', dtype={'seq': str, 'frame': str})  
    gt['label'] = gt['label'].apply(lambda x: ast.literal_eval(x) if pd.notna(x) else None)  
  
    # plots one frame  
    seq = '000'  
    frame = '6'  
    plot_frame(pred, gt, seq, frame, RESULTS_PATH)  
  
    # plot all frames from all sequences  
    seqs = gt['seq'].unique()  
  
    # plots whole sequence  
    for seq in seqs:  
        print("processing seq %s" % seq)  
        plot_sequence(pred, gt, seq, RESULTS_PATH)
```

Select which sequence or frame you want to inspect.

# The Baseline Solution

## evaluate.py

```
if __name__=="__main__":  
    mAP, AP = compute_mAP_from_files("predictions.csv", "/home/master/dataset/test/labels.csv")  
    print("mAP:{:.4f}".format(mAP))  
    for ap_metric, iou in zip(AP, np.arange(0.5, 1, 0.05)):  
        print("\tmAP at IoU level [{:.2f}]: {:.4f}".format(iou, ap_metric))
```

Compute scores by calling the metrics function.



## Ideas / Suggestions:

- Tweak some model (hyper)parameters (anchors, backbones, data augmentation, data splitting);
- Find ways to deal with the big variability of visibility conditions underwater;
- Analyse sequences instead of individual images (use temporal information/context).

# Start Exploring

# Next Steps (Recommended)

## Information

- Go to <https://github.com/visum-summer-school/visum-competition2020>
- Read the visum\_project\_FAQ.pdf file

## Team Registration

- Organize into a group of **three** and register your team.
- If you do not have a team and want to participate, send us an e-mail and we will find you a team!

## First result

- Set up a valid test.py and get your name on the leaderboard! You may use the provided baseline for this.